

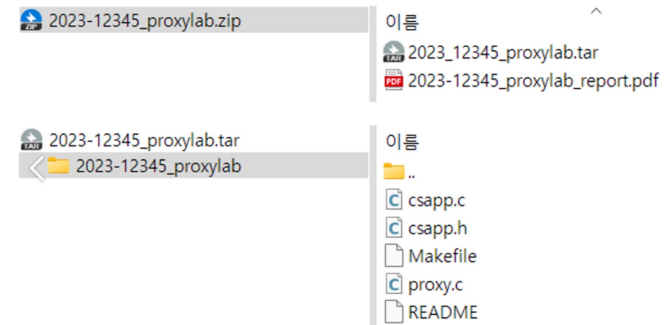
System Programming Lab #5

2023-05-23

sp-tas

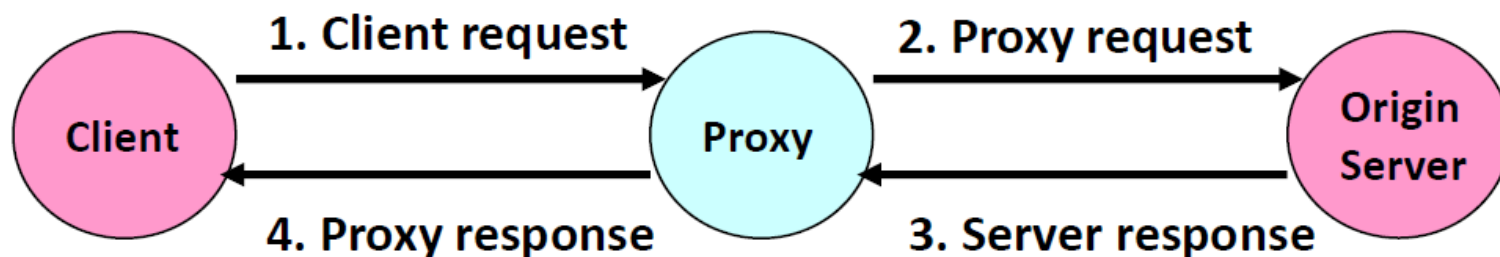
Lab Assignment #5 : Proxy Lab

- Download skeleton code & pdf from eTL
proxylab-handout.tar, proxylab-handout.pdf
- Hand In
 - 구현 디렉토리 압축파일: 학번-proxylab.tar
 - Upload your files eTL
 - 압축파일 양식 : [학번]_proxylab.zip
 - Ex) 2023-12345_proxylab.zip
 - A zip file should include
 - (1) a tarball of your implementation directory (2) report
 - tarball 양식 : [학번]_proxylab.tar eg) 2023-12345_proxylab.tar
 - Report 양식 : [학번]_proxylab_report.pdf
 - Please, **READ** the Hand-out and Lab material thoroughly!
- Assigned : May 23th
- Deadline : June 20th, 23:59:59 (**No late submission**)



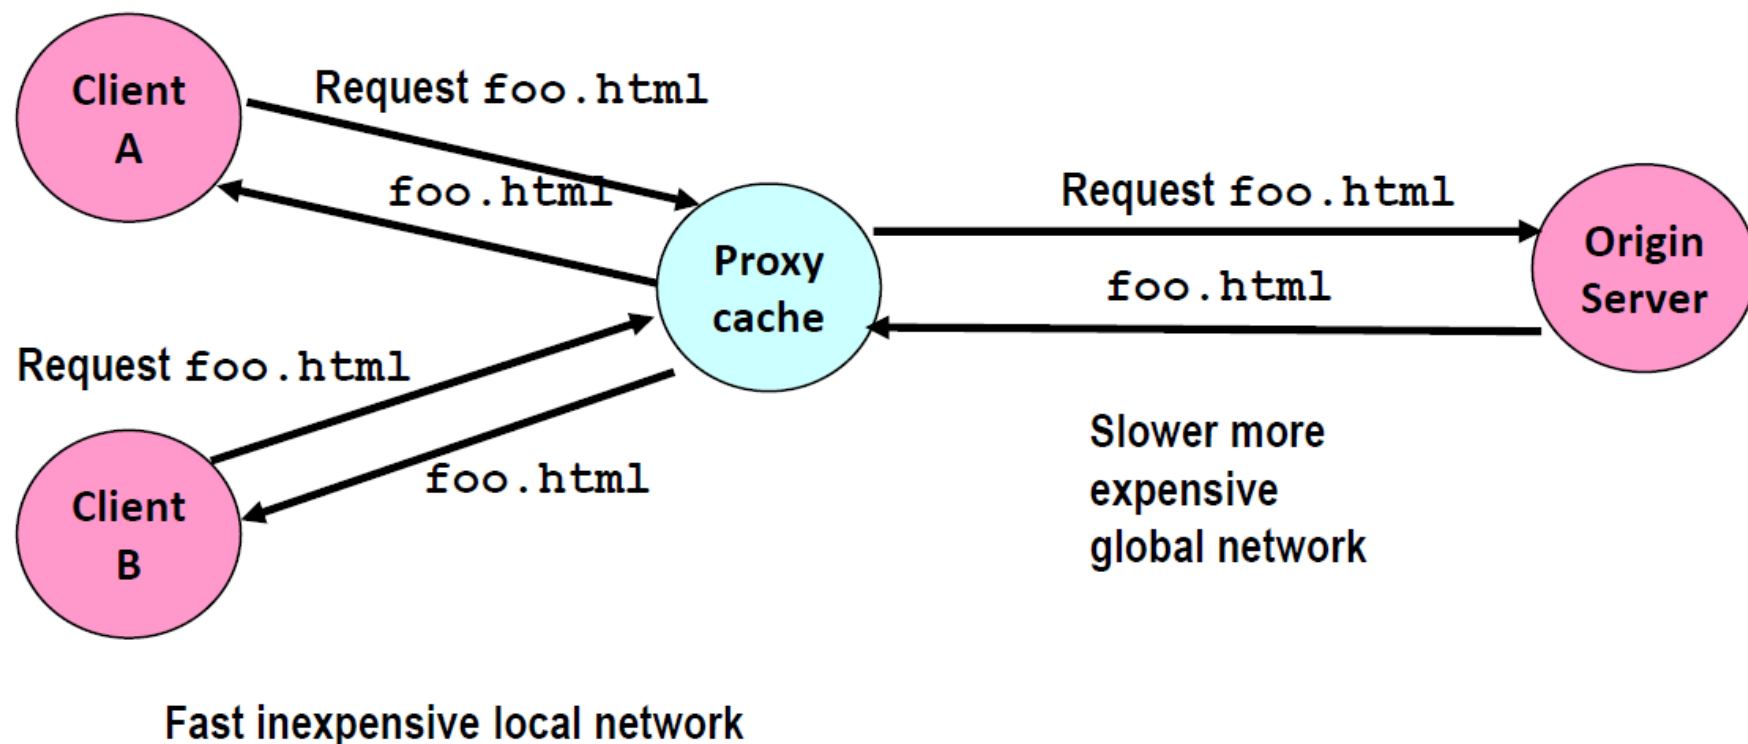
Proxies

- A **proxy** is an intermediary between a client and an **origin server**
 - To the client, the proxy acts like a server
 - To the server, the proxy acts like a client



Why Proxies?

- Can perform useful functions as requests and responses pass by
 - Examples: Caching, logging, anonymization, filtering, transcoding



How the Web Really Works

- In reality, a single HTML page today may depend on 10s or 100s of support files (images, stylesheets, scripts, etc.)
- Builds a good argument for concurrent servers
 - Just to load a single modern webpage, the client would have to wait for 10s of back-to-back request
 - I/O is likely slower than processing, so back
- Caching is simpler if done in pieces rather than whole page
 - If only part of the page changes, no need to fetch old parts again
 - Each object (image, stylesheet, script) already has a unique URL that can be used as a key

You will implement

- Write a simple HTTP proxy that caches web objects
- Part 1: Implementing a sequential web Proxy
 - **Basic HTTP operation & socket programming**
 - set up the proxy to accept incoming connections
 - read and parse requests
 - forward requests to web servers
 - read the servers' responses
 - forward those responses to the corresponding clients
- Part 2: Dealing with multiple concurrent requests
 - upgrade your proxy to deal with multiple **concurrent** connections
 - multi-threading
- Part 3: Caching web objects
 - add caching to your proxy using a simple main memory cache of recently accessed web content
 - cache individual objects, not the whole page
 - **Use an LRU eviction policy**
 - your caching system must allow for concurrent reads while maintaining consistency

Guide to start your implementation

- `int main(int argc, char *argv[])`
 - initialize everything such as data structure
 - checking port number
 - establish listening requests
 - when a client connects, spawn a new thread to handle it

```
1  #include <stdio.h>
2
3  /* Recommended max cache and object sizes */
4  #define MAX_CACHE_SIZE 1049000
5  #define MAX_OBJECT_SIZE 102400
6
7  /* You won't lose style points for including this long line in your code */
8  static const char *user_agent_hdr = "User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:10.0.3) Gecko/20120305 Firefox/10.0.3\r\n";
9
10 int main()
11 {
12     printf("%s", user_agent_hdr);
13     return 0;
14 }
15
```

Use csapp.[ch] functions

- Also, csapp.[ch] codes are included! yeah!

```
/* Sockets interface wrappers */
int Socket(int domain, int type, int protocol);
void Setsockopt(int s, int level, int optname, const void *optval, int optlen);
void Bind(int sockfd, struct sockaddr *my_addr, int addrlen);
void Listen(int s, int backlog);
int Accept(int s, struct sockaddr *addr, socklen_t *addrlen);
void Connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
```

```
/* Protocol independent wrappers */
void Getaddrinfo(const char *node, const char *service,
                 const struct addrinfo *hints, struct addrinfo **res);
void Getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host,
                 size_t hostlen, char *serv, size_t servlen, int flags);
void Freeaddrinfo(struct addrinfo *res);
void Inet_ntop(int af, const void *src, char *dst, socklen_t size);
void Inet_pton(int af, const char *src, void *dst);
```

```
/* DNS wrappers */
struct hostent *Gethostbyname(const char *name);
struct hostent *Gethostbyaddr(const char *addr, int len, int type);
```

```
/* Pthreads thread control wrappers */
void Pthread_create(pthread_t *tidp, pthread_attr_t *attrp,
                   void * (*routine)(void *), void *argp);
void Pthread_join(pthread_t tid, void **thread_return);
void Pthread_cancel(pthread_t tid);
void Pthread_detach(pthread_t tid);
void Pthread_exit(void *retval);
pthread_t Pthread_self(void);
void Pthread_once(pthread_once_t *once_control, void (*init_function)());
```

```
/* Rio (Robust I/O) package */
ssize_t rio_readn(int fd, void *usrbuf, size_t n);
ssize_t rio_writen(int fd, void *usrbuf, size_t n);
void rio_readinitb(rio_t *rp, int fd);
ssize_t rio_readnb(rio_t *rp, void *usrbuf, size_t n);
ssize_t rio_readlineb(rio_t *rp, void *usrbuf, size_t maxlen);
```

```
/* Wrappers for Rio package */
ssize_t Rio_readn(int fd, void *usrbuf, size_t n);
void Rio_writen(int fd, void *usrbuf, size_t n);
void Rio_readinitb(rio_t *rp, int fd);
ssize_t Rio_readnb(rio_t *rp, void *usrbuf, size_t n);
ssize_t Rio_readlineb(rio_t *rp, void *usrbuf, size_t maxlen);
```

```
/* Reentrant protocol-independent client/server helpers */
int open_clientfd(char *hostname, char *port);
int open_listenfd(char *port);
```

```
/* Wrappers for reentrant protocol-independent client/server helpers */
int Open_clientfd(char *hostname, char *port);
int Open_listenfd(char *port);
```


Checking Your Work

- Auto grader
 - `./driver.sh` will run the tests:
 - Ability to pull basic web pages from a server
 - Handle a (concurrent) request while another request is still pending
 - Fetch a web page again from your cache after the server has been stopped
 - This should help answer the question:
“Is this what my proxy is supposed to do?”
 - Please don’t use this grader to definitively test your proxy; there are many things not tested here
 - If you have a permission problem, use the command below
 - `chmod +x *`

Checking Your Work

- Get your port
 - Generate a port for yourself with `./port-for-user.pl [sp ID]`
 - Generate more ports for web servers and such with `./free-port.sh`
- Run proxy server
 - `./proxy port_id &` (background)
- Request to server
 - (without proxy) `>> curl http://www.sk.co.kr`
 - (with proxy) `>> curl --proxy localhost:port_id http://www.sk.co.kr`

Evaluation

- Total Score: 100 points
- Basic Correctness (40 points)
 - basic proxy operation (auto graded)
- Concurrency (15 points)
 - handling concurrent requests (auto graded)
- Cache (15 points)
 - working cache (auto graded)
- Real Pages (20 points)
 - correctly serving the real pages
- Report (10 points)
 - describes the goal of proxy lab and how to implement for each part
 - what you learn in this lab
 - what was difficult, surprising, and so on

Don't forget!

Last year's FAQ

- Q1. Do I need to implement GET request only?
 - A1. Yes. Other requests (e.g., POST) are optional.
- Q2. Do I have to consider chunked responses?
 - A2. No, this is also optional.
- Q3. May I assume that the URI of a GET request is an absolute path? (e.g., `http://example.com/index.html`)
 - A3. Yes, relative paths (e.g., `/index.html`) are not tested.
- Q4. Which size is used for calculating the cache size?
 - A4. The size of a response message from the server is used.

Last year's FAQ

- Q5. Does the response of the same request can be changed?
 - A5. In our evaluation, the response will be the same.
- Q6. Timeout bug in `driver.sh`?
 - A6. Check if `python` and `netstat` are installed properly.
- Q7. Does the ordering of the header fields affect the response?
 - A7. No, it does not affect the response of a request.

Fin.

- Questions
 - eTL Q&A Board
- Read the handout thoroughly & start early!
- This is our last LAB session
- The QnA session will 6/13